

# hw3\_cnn\_for\_ieng-Alexnet

May 14, 2018

```
In [7]: import math
import timeit
import numpy as np
import matplotlib.pyplot as plt

import torch
import torchvision
import torch.nn as nn
import torch.optim as optim
import torch.nn.init as init
import torch.nn.functional as F
from torch.autograd import Variable
import torch.backends.cudnn as cudnn
import torchvision.transforms as transforms
from torch.utils.data.sampler import SubsetRandomSampler

In [8]: # Process data
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

# Load Train data and split into train set and valid set
num_train = 50000
indices = list(range(num_train))
train_index, valid_index = indices[10000:], indices[:10000]
train_sampler=SubsetRandomSampler(train_index)
valid_sampler=SubsetRandomSampler(valid_index)
train_dataset = torchvision.datasets.CIFAR10(root='/datasets/CIFAR-10', train=True,
                                             download=False, transform=transform)
valid_dataset = torchvision.datasets.CIFAR10(root='/datasets/CIFAR-10', train=True,
                                             download=False, transform=transform)
train_loader = torch.utils.data.DataLoader(train_dataset,
                                             batch_size=32, sampler=train_sampler,
                                             num_workers=4, pin_memory=True)
valid_loader = torch.utils.data.DataLoader(valid_dataset,
                                             batch_size=10, sampler=valid_sampler,
                                             num_workers=4, pin_memory=True)
```

```

# Load test data
testset = torchvision.datasets.CIFAR10(root='/datasets/CIFAR-10', train=False,
                                       download=False, transform=transform)
test_loader = torch.utils.data.DataLoader(testset,
                                          batch_size=10, pin_memory=True,
                                          shuffle=True, num_workers=4)

# Define Class
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

In [9]: # Save data in list for plot
save_train_loss=[]
save_valid_loss=[]
save_test_loss=[]
save_train_acc=[]
save_valid_acc=[]
save_test_acc=[]

In [13]: def main(iteration = 10, gpu_usage = False, opti = 'default',
                  Xavier = False, batch_norm = False):
    # detect GPU
    if gpu_usage:
        use_cuda = torch.cuda.is_available()
    else:
        use_cuda = False
    # define network
    class Net(nn.Module):
        def __init__(self):
            super(Net, self).__init__()
            if batch_norm:
                # batch normalization
                self.conv1 = nn.Conv2d(3, 96, 3, stride=1, padding=1)
                self.conv1_bn = nn.BatchNorm2d(96)
                self.pool1 = nn.MaxPool2d(2, 2)
                self.conv2 = nn.Conv2d(96, 256, 3, stride=1, padding=1)
                self.conv2_bn = nn.BatchNorm2d(256)
                self.pool2 = nn.MaxPool2d(2, 2)
                self.conv3 = nn.Conv2d(256, 384, 3, stride=1, padding=1)
                self.conv3_bn = nn.BatchNorm2d(384)
                self.conv4 = nn.Conv2d(384, 384, 3, stride=1, padding=1)
                self.conv4_bn = nn.BatchNorm2d(384)
                self.conv5 = nn.Conv2d(384, 256, 3, stride=1, padding=1)
                self.conv5_bn = nn.BatchNorm2d(256)
                self.pool5 = nn.MaxPool2d(2, 2)
                self.fc1 = nn.Linear(256 * 4 * 4, 4096)
                self.drop1 = nn.Dropout2d(0.2)

```

```

        self.fc2 = nn.Linear(4096, 4096)
        self.drop2 = nn.Dropout2d(0.2)
        self.fc3 = nn.Linear(4096, 10)
    else:
        # no batch normalization
        self.conv1 = nn.Conv2d(3, 96, 3, stride=1, padding=1)
        self.pool1 = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(96, 256, 3, stride=1, padding=1)
        self.pool2 = nn.MaxPool2d(2, 2)
        self.conv3 = nn.Conv2d(256, 384, 3, stride=1, padding=1)
        self.conv4 = nn.Conv2d(384, 384, 3, stride=1, padding=1)
        self.conv5 = nn.Conv2d(384, 256, 3, stride=1, padding=1)
        self.pool5 = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(256 * 4 * 4, 4096)
        self.drop1 = nn.Dropout2d(0.2)
        self.fc2 = nn.Linear(4096, 4096)
        self.drop2 = nn.Dropout2d(0.2)
        self.fc3 = nn.Linear(4096, 10)

    # Xavier init, need to check formula
    # Xavier Initialize for the whole network
    if Xavier:
        for m in self.modules():
            if isinstance(m, nn.Conv2d):
                n = m.kernel_size[0] * m.kernel_size[1] * m.out_channels
                m.weight.data.normal_(0, math.sqrt(2. / n))
            elif isinstance(m, nn.BatchNorm2d):
                m.weight.data.fill_(1)
                m.bias.data.zero_()

def forward(self, x):
    if batch_norm:
        # batch normalization
        x = self.pool1(F.relu(self.conv1_bn(self.conv1(x))))
        x = self.pool2(F.relu(self.conv2_bn(self.conv2(x))))
        x = F.relu(self.conv3_bn(self.conv3(x)))
        x = F.relu(self.conv4_bn(self.conv4(x)))
        x = self.pool5(F.relu(self.conv5_bn(self.conv5(x))))
        x = x.view(-1, 256 * 4 * 4)
        x = self.drop1(F.relu(self.fc1(x)))
        x = self.drop2(F.relu(self.fc2(x)))
        x = self.fc3(x)
    else:
        # no batch normalization
        x = self.pool1(F.relu(self.conv1(x)))
        x = self.pool2(F.relu(self.conv2(x)))
        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))

```

```

        x = self.pool5(F.relu(self.conv5(x)))
        x = x.view(-1, 256 * 4 * 4)
        x = self.drop1(F.relu(self.fc1(x)))
        x = self.drop2(F.relu(self.fc2(x)))
        x = self.fc3(x)
    return x

net = Net()
# net.apply(weights_init)

## using GPU
if use_cuda:
    net.cuda()
    net = torch.nn.DataParallel(net, device_ids=range(torch.cuda.device_count()))
    cudnn.benchmark = True

## loss function
criterion = nn.CrossEntropyLoss()
if (opti == "Adam"):
    optimizer = optim.Adam(net.parameters(), lr = 0.001)
else:
    optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.9, nesterov=True)

## starting training testinig validating
start = timeit.default_timer()
print('epoch, mini-batch\tttrain_loss\ttvalid_loss\tttest_loss\tttrain_acc\ttvalid_acc\tttest_acc')
for epoch in range(iteration): # loop over the dataset multiple times
    train_running_loss = 0.0
    for i, data in enumerate(train_loader, 0):
        # get the inputs
        inputs, labels = data
        # detect GPU
        if use_cuda:
            inputs, labels = inputs.cuda(), labels.cuda()
        # wrap them in Variable
        inputs, labels = Variable(inputs), Variable(labels)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
    train_running_loss += loss.data[0]

```

```

if i % 501 == 500:    # print every 1 mini-batches
    train_loss = train_running_loss / 500
    #print('train_loss = '+str(train_loss))
    train_running_loss = 0.0
    # Train
    correct = 0
    total = 0
    for data in train_loader:
        images, labels = data
        # detect GPU
        if use_cuda:
            images, labels = images.cuda(), labels.cuda()
            outputs = net(Variable(images))
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum()
    train_acc = 100 * correct / total
    #print('train_acc = '+str(train_acc))
    # Valid
    correct = 0
    total = 0
    valid_loss = 0
    for data in valid_loader:
        images, labels = data
        # detect GPU
        if use_cuda:
            images, labels = images.cuda(), labels.cuda()
            v_images, v_labels = Variable(images), Variable(labels)
            outputs = net(v_images)
            valid_loss += criterion(outputs, v_labels).data[0]
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum()
    valid_loss /= 1000
    valid_acc = 100 * correct / total
    # Test
    correct = 0
    total = 0
    test_loss = 0
    for data in test_loader:
        images, labels = data
        # detect GPU
        if use_cuda:
            images, labels = images.cuda(), labels.cuda()
            v_images, v_labels = Variable(images), Variable(labels)
            outputs = net(v_images)
            test_loss += criterion(outputs, v_labels).data[0]

```

```

        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum()
    test_loss /= 1000
    test_acc = 100 * correct / total
    # Save data for plot
    save_train_loss.append(train_loss)
    save_valid_loss.append(valid_loss)
    save_test_loss.append(test_loss)
    save_train_acc.append(train_acc)
    save_valid_acc.append(valid_acc)
    save_test_acc.append(test_acc)
    print('[%d, %5d]\t\t%.3f\t\t%.3f\t\t%.3f\t\t%.3f%%\t\t%.3f%%\t\t%.3f%%'
          % (iteration, test_loss, test_acc, train_loss, train_acc, valid_loss, valid_acc))
    print('Finished Training')
    stop = timeit.default_timer()
    print('Time it takes: %.3f second' % (stop - start))

```

```

In [14]: # iteration: howmany iteration you want. default = 10
         # gpu_usage: use GPU or not. default = false
         # opti: type of optimizer. default: SGD
         main(iteration = 100, gpu_usage = True, opti = 'SGD', Xavier = True, batch_norm = True)

```

epoch, mini-batch	train_loss	valid_loss	test_loss	train_acc	valid_acc
[1, 501]	1.542	1.334	1.338	54.227%	54.227%
[1, 1002]	1.194	1.157	1.165	62.138%	62.138%
[2, 501]	0.953	0.913	0.923	71.957%	71.957%
[2, 1002]	0.864	0.889	0.906	73.793%	73.793%
[3, 501]	0.714	0.852	0.897	76.500%	76.500%
[3, 1002]	0.714	0.771	0.810	79.815%	79.815%
[4, 501]	0.578	0.740	0.747	81.880%	81.880%
[4, 1002]	0.590	0.747	0.767	82.448%	82.448%
[5, 501]	0.471	0.758	0.766	84.290%	84.290%
[5, 1002]	0.489	0.664	0.676	86.972%	86.972%
[6, 501]	0.361	0.729	0.717	87.892%	87.892%
[6, 1002]	0.408	0.679	0.671	89.278%	89.278%
[7, 501]	0.286	0.695	0.727	91.263%	91.263%
[7, 1002]	0.320	0.698	0.719	91.302%	91.302%
[8, 501]	0.226	0.715	0.746	92.463%	92.463%
[8, 1002]	0.250	0.710	0.727	94.740%	94.740%
[9, 501]	0.174	0.756	0.783	94.203%	94.203%
[9, 1002]	0.203	0.702	0.737	95.647%	95.647%
[10, 501]	0.127	0.768	0.822	95.422%	95.422%
[10, 1002]	0.144	0.716	0.734	96.207%	96.207%
[11, 501]	0.107	0.806	0.836	96.500%	96.500%
[11, 1002]	0.117	0.776	0.821	96.858%	96.858%
[12, 501]	0.081	0.781	0.794	98.020%	98.020%
[12, 1002]	0.095	0.801	0.857	97.670%	97.670%
[13, 501]	0.064	0.833	0.848	97.525%	97.525%

[13, 1002]	0.082	0.826	0.844	98.078
[14, 501]	0.051	0.886	0.950	98.040
[14, 1002]	0.061	0.923	0.952	97.862
[15, 501]	0.048	0.867	0.902	98.640
[15, 1002]	0.062	0.850	0.868	98.070
[16, 501]	0.045	0.843	0.863	99.188
[16, 1002]	0.038	0.838	0.879	99.203
[17, 501]	0.029	0.973	1.032	98.835
[17, 1002]	0.036	0.943	1.004	98.618
[18, 501]	0.032	0.867	0.925	99.237
[18, 1002]	0.038	0.888	0.951	99.040
[19, 501]	0.023	0.871	0.901	99.595
[19, 1002]	0.016	0.920	0.945	99.545
[20, 501]	0.022	0.921	0.977	99.297
[20, 1002]	0.019	0.895	0.958	99.465
[21, 501]	0.015	0.920	0.976	99.690
[21, 1002]	0.027	0.879	0.960	99.388
[22, 501]	0.027	0.882	0.931	99.595
[22, 1002]	0.014	0.920	1.024	99.740
[23, 501]	0.015	0.972	1.033	99.655
[23, 1002]	0.019	0.927	0.983	99.540
[24, 501]	0.023	0.978	1.039	99.218
[24, 1002]	0.018	0.941	1.025	99.535
[25, 501]	0.019	0.948	1.007	99.493
[25, 1002]	0.022	1.013	1.062	99.320
[26, 501]	0.014	0.933	0.975	99.713
[26, 1002]	0.013	0.996	1.061	99.510
[27, 501]	0.015	0.934	0.966	99.677
[27, 1002]	0.014	1.041	1.074	99.493
[28, 501]	0.018	0.983	1.041	99.547
[28, 1002]	0.012	0.961	1.000	99.740
[29, 501]	0.015	0.984	1.041	99.570
[29, 1002]	0.017	0.943	1.015	99.630
[30, 501]	0.012	0.981	1.062	99.767
[30, 1002]	0.012	0.950	1.011	99.805
[31, 501]	0.013	0.986	1.008	99.700
[31, 1002]	0.016	0.964	1.030	99.600
[32, 501]	0.011	0.992	1.046	99.567
[32, 1002]	0.008	1.015	1.073	99.755
[33, 501]	0.013	1.016	1.049	99.588
[33, 1002]	0.012	0.990	1.062	99.627
[34, 501]	0.012	0.994	1.053	99.662
[34, 1002]	0.013	1.020	1.076	99.608
[35, 501]	0.010	1.022	1.025	99.685
[35, 1002]	0.010	1.007	1.038	99.797
[36, 501]	0.006	1.002	1.033	99.892
[36, 1002]	0.006	1.044	1.065	99.793
[37, 501]	0.005	1.028	1.092	99.778

[37, 1002]	0.005	1.061	1.103	99.853
[38, 501]	0.007	1.039	1.104	99.820
[38, 1002]	0.008	1.047	1.101	99.832
[39, 501]	0.005	1.118	1.212	99.665
[39, 1002]	0.009	1.038	1.088	99.850
[40, 501]	0.007	1.182	1.254	99.233
[40, 1002]	0.011	1.086	1.160	99.800

Process Process-1127:

Process Process-1126:

Process Process-1128:

Process Process-1125:

KeyboardInterruptTraceback (most recent call last)

```

<ipython-input-14-9eb1a9fc7a7f> in <module>()
      2 # gpu_usage: use GPU or not. default = false
      3 # opti: type of optimizer. default: SGD
----> 4 main(iteration = 100, gpu_usage = True, opti = 'SGD', Xavier = True, batch_norm = Tr

<ipython-input-13-8f4b7d1c7234> in main(iteration, gpu_usage, opti, Xavier, batch_norm)
    118         outputs = net(inputs)
    119         loss = criterion(outputs, labels)
--> 120         loss.backward()
    121         optimizer.step()
    122

/opt/conda/lib/python3.6/site-packages/torch/autograd/variable.py in backward(self, grad
    165         Variable.
    166         """
--> 167         torch.autograd.backward(self, gradient, retain_graph, create_graph, retain_v
    168
    169         def register_hook(self, hook):

/opt/conda/lib/python3.6/site-packages/torch/autograd/__init__.py in backward(variables,
    97
    98     Variable._execution_engine.run_backward(
--> 99         variables, grad_variables, retain_graph)
    100
    101

```



KeyboardInterrupt:

Traceback (most recent call last):

Traceback (most recent call last):

Traceback (most recent call last):

Traceback (most recent call last):

File "/opt/conda/lib/python3.6/multiprocessing/process.py", line 258, in \_bootstrap  
self.run()

File "/opt/conda/lib/python3.6/multiprocessing/process.py", line 258, in \_bootstrap  
self.run()

File "/opt/conda/lib/python3.6/multiprocessing/process.py", line 258, in \_bootstrap  
self.run()

File "/opt/conda/lib/python3.6/multiprocessing/process.py", line 258, in \_bootstrap  
self.run()

File "/opt/conda/lib/python3.6/multiprocessing/process.py", line 93, in run  
self.\_target(\*self.\_args, \*\*self.\_kwargs)

File "/opt/conda/lib/python3.6/multiprocessing/process.py", line 93, in run  
self.\_target(\*self.\_args, \*\*self.\_kwargs)

File "/opt/conda/lib/python3.6/multiprocessing/process.py", line 93, in run  
self.\_target(\*self.\_args, \*\*self.\_kwargs)

File "/opt/conda/lib/python3.6/multiprocessing/process.py", line 93, in run  
self.\_target(\*self.\_args, \*\*self.\_kwargs)

File "/opt/conda/lib/python3.6/site-packages/torch/utils/data/dataloader.py", line 50, in \_worker\_loop  
r = index\_queue.get()

File "/opt/conda/lib/python3.6/site-packages/torch/utils/data/dataloader.py", line 50, in \_worker\_loop  
r = index\_queue.get()

File "/opt/conda/lib/python3.6/site-packages/torch/utils/data/dataloader.py", line 50, in \_worker\_loop  
r = index\_queue.get()

File "/opt/conda/lib/python3.6/site-packages/torch/utils/data/dataloader.py", line 55, in \_worker\_loop  
samples = collate\_fn([dataset[i] for i in batch\_indices])

File "/opt/conda/lib/python3.6/multiprocessing/queues.py", line 334, in get  
with self.\_rlock:

File "/opt/conda/lib/python3.6/multiprocessing/queues.py", line 334, in get  
with self.\_rlock:

File "/opt/conda/lib/python3.6/multiprocessing/queues.py", line 335, in get  
res = self.\_reader.recv\_bytes()

File "/opt/conda/lib/python3.6/site-packages/torch/utils/data/dataloader.py", line 55, in \_worker\_loop  
samples = collate\_fn([dataset[i] for i in batch\_indices])

File "/opt/conda/lib/python3.6/multiprocessing/synchronize.py", line 96, in \_\_enter\_\_  
return self.\_semlock.\_\_enter\_\_()

File "/opt/conda/lib/python3.6/multiprocessing/synchronize.py", line 96, in \_\_enter\_\_  
return self.\_semlock.\_\_enter\_\_()

File "/opt/conda/lib/python3.6/multiprocessing/connection.py", line 216, in recv\_bytes  
buf = self.\_recv\_bytes(maxlength)

File "/opt/conda/lib/python3.6/site-packages/torchvision-0.2.0-py3.6.egg/torchvision/datasets/

```

    img = Image.fromarray(img)
KeyboardInterrupt
KeyboardInterrupt
File "/opt/conda/lib/python3.6/multiprocessing/connection.py", line 407, in _recv_bytes
    buf = self._recv(4)
File "/opt/conda/lib/python3.6/site-packages/PIL/Image.py", line 2446, in fromarray
    obj = obj.tobytes()
File "/opt/conda/lib/python3.6/multiprocessing/connection.py", line 379, in _recv
    chunk = read(handle, remaining)
KeyboardInterrupt
KeyboardInterrupt

```

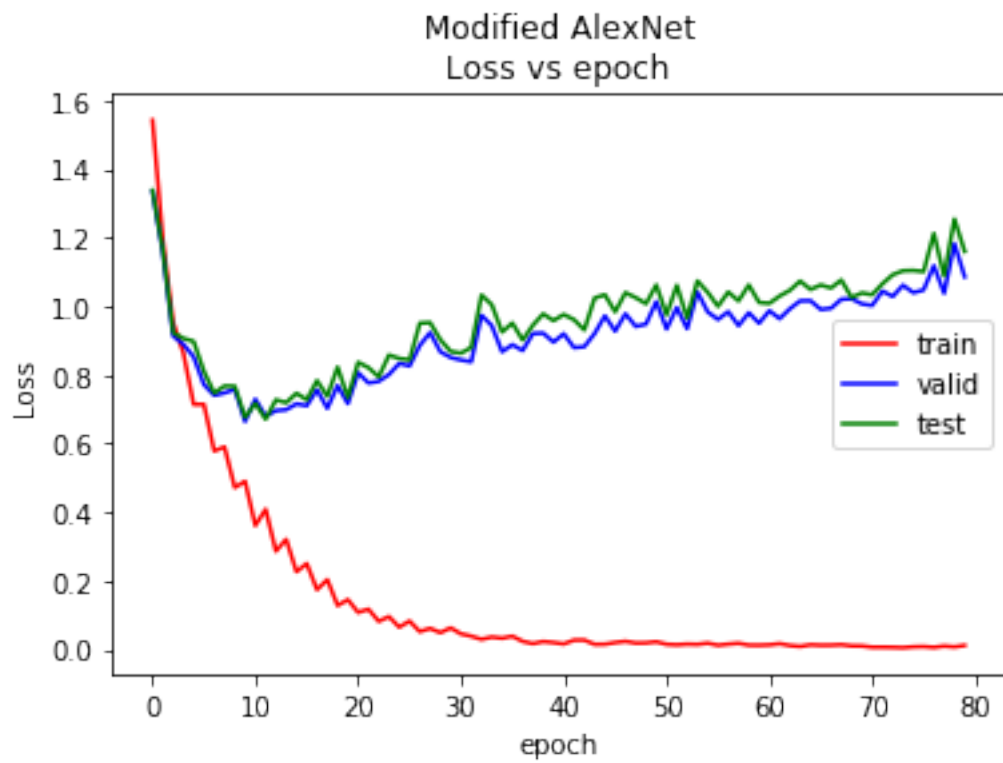
```

In [18]: def plot_accuracy(trainacc,validacc,testacc):
    plt.figure()
    plt.title('Modified AlexNet\nAccuracy vs epoch')
    plt.ylabel('Accuracy')
    plt.xlabel('epoch')
    plt.plot(trainacc,'red')
    plt.plot(validacc,'blue')
    plt.plot(testacc,'green')
    plt.legend(['train','valid','test'])
    plt.savefig('HW3_accuracy_model_3_32_batch_001')
    plt.show()

    def plot_loss(trainloss,validloss,testloss):
        plt.figure()
        plt.title('Modified AlexNet\nLoss vs epoch')
        plt.ylabel('Loss')
        plt.xlabel('epoch')
        plt.plot(trainloss,'red')
        plt.plot(validloss,'blue')
        plt.plot(testloss,'green')
        plt.legend(['train','valid','test'])
        plt.savefig('HW3_loss_model_3_32_batch_001')
        plt.show()

In [19]: plot_loss(save_train_loss,save_valid_loss,save_test_loss)

```



```
In [20]: plot_accuracy(save_train_acc,save_valid_acc,save_test_acc)
```

